

superstore_case

December 7, 2025

```
[1]: import pandas as pd
import numpy as np
file_path ="C:/Users/anush/Downloads/superstore_eda_v1-1724655032.csv"
df = pd.read_csv(file_path)
```

```
[2]: df.shape
```

```
[2]: (10014, 21)
```

```
[3]: df.head(5)
```

```
[3]:   Row ID      Order ID Order Date Ship Date      Ship Mode Customer ID \
0      1 CA-2016-152156  08/11/16  11/11/16 Second Class CG-12520
1      2 CA-2016-152156  08/11/16  11/11/16 Second Class CG-12520
2      3 CA-2016-138688  12/06/16  16/06/16 Second Class DV-13045
3      4 US-2015-108966  11/10/15  18/10/15 Standard Class SO-20335
4      5 US-2015-108966  11/10/15  18/10/15 Standard Class SO-20335
```

```
      Customer Name    Segment      Country          City ... \
0      Claire Gute    Consumer United States Henderson ...
1      Claire Gute    Consumer United States Henderson ...
2 Darrin Van Huff    Corporate United States Los Angeles ...
3 Sean O'Donnell    Consumer United States Fort Lauderdale ...
4 Sean O'Donnell    Consumer United States Fort Lauderdale ...
```

```
  Postal Code Region      Product ID      Category Sub-Category \
0      42420   South FUR-BO-10001798 Furniture Bookcases
1      42420   South FUR-CH-10000454 Furniture Chairs
2     90036    West OFF-LA-10000240 Office Supplies Labels
3     33311   South FUR-TA-10000577 Furniture Tables
4     33311   South OFF-ST-10000760 Office Supplies Storage
```

```
      Product Name Sales Price  Quantity \
0      Bush Somerset Collection Bookcase  261.9600    2.0
1 Hon Deluxe Fabric Upholstered Stacking Chairs,...  731.9400    3.0
2 Self-Adhesive Address Labels for Typewriters b...  14.6200    2.0
3 Bretford CR4500 Series Slim Rectangular Table  957.5775    5.0
4 Eldon Fold 'N Roll Cart System  22.3680    2.0
```

```
Discount      Profit
0       0.00    41.9136
1       0.00   219.5820
2       0.00    6.8714
3       0.45 -383.0310
4       0.20    2.5164
```

[5 rows x 21 columns]

[4]: #Handling Duplicates:

```
df = df.apply(lambda x: x.str.strip() if x.dtype == "object" else x)
df = df.apply(lambda x: x.str.lower() if x.dtype == "object" else x)

print("Duplicate rows before:", df.duplicated().sum())

rows_before = df.shape[0]

df = df.drop_duplicates()

rows_after = df.shape[0]

print("\nAfter cleaning:")
print("Rows before:", rows_before)
print("Rows after:", rows_after)
print("Rows removed:", rows_before - rows_after)
```

Duplicate rows before: 17

After cleaning:

Rows before: 10014

Rows after: 9997

Rows removed: 17

[5]: df.dtypes

```
Row ID          int64
Order ID        object
Order Date      object
Ship Date       object
Ship Mode       object
Customer ID     object
Customer Name   object
Segment         object
Country         object
City            object
State           object
```

```
Postal Code      int64
Region          object
Product ID     object
Category        object
Sub-Category   object
Product Name   object
Sales Price    float64
Quantity       float64
Discount       float64
Profit         float64
dtype: object
```

```
[6]: df.describe()
```

```
[6]:      Row ID  Postal Code  Sales Price  Quantity  Discount \
count  9997.000000  9997.000000  9997.000000  9979.000000  9997.000000
mean   4996.992198  55192.055117   229.798602   3.790861   0.156216
std    2885.688679  32062.477713   623.164427   2.226648   0.206422
min    1.000000   1040.000000  -31.500000   1.000000   0.000000
25%   2498.000000  23223.000000   17.240000   2.000000   0.000000
50%   4997.000000  56560.000000   54.480000   3.000000   0.200000
75%   7496.000000  90008.000000  209.940000   5.000000   0.200000
max   9994.000000  99301.000000  22638.480000  14.000000   0.800000

      Profit
count  9997.000000
mean   28.650525
std    234.225264
min   -6599.978000
25%   1.731000
50%   8.662000
75%   29.364000
max   8399.976000
```

```
[7]: # Correct date format:-
```

```
df["Ship Date"] = pd.to_datetime(df["Ship Date"], errors='coerce')
df["Order Date"] = pd.to_datetime(df["Order Date"], errors='coerce')
print(df[['Order Date', 'Ship Date']].dtypes)
```

```
Order Date      datetime64[ns]
Ship Date      datetime64[ns]
dtype: object
```

```
C:\Users\anush\AppData\Local\Temp\ipykernel_22460\3538689358.py:3: UserWarning:
Could not infer format, so each element will be parsed individually, falling
back to `dateutil`. To ensure parsing is consistent and as-expected, please
specify a format.
```

```

df["Ship Date"] = pd.to_datetime(df["Ship Date"], errors='coerce')
C:\Users\anush\AppData\Local\Temp\ipykernel_22460\3538689358.py:4: UserWarning:
Could not infer format, so each element will be parsed individually, falling
back to `dateutil`. To ensure parsing is consistent and as-expected, please
specify a format.

df["Order Date"] = pd.to_datetime(df["Order Date"], errors='coerce')

```

[8]: # Extract the year from the Order ID and compare it with the year in Order Date.
 ↪ Correct any inconsistencies.

```

df['Order_Year_From_ID'] = df["Order ID"].str.split('-').str[1]
df['Order_Year_From_ID'] = pd.to_numeric(df['Order_Year_From_ID'], errors = "coerce")
df['Order_Year_From_Date'] = df['Order Date'].dt.year

mismatch_mask = df['Order_Year_From_ID'] != df['Order_Year_From_Date']
print("Number of mismatched rows:", mismatch_mask.sum())

df.loc[mismatch_mask, 'Order Date'] = df.loc[mismatch_mask].apply(lambda row: row['Order Date'].replace(year=row['Order_Year_From_ID']))
    if pd.notnull(row['Order Date']) else row['Order Date'], axis=1)

```

Number of mismatched rows: 40

[9]: # Negative values in Sales price

```

print(df[df["Sales Price"] < 0].shape[0])
df= df[df["Sales Price"] >= 0]

```

4

[10]: df.describe()

	Row ID	Order Date	\		
count	9993.000000	9993			
mean	4996.022716	2016-04-11 14:27:37.880516352			
min	1.000000	2014-01-02 00:00:00			
25%	2498.000000	2015-05-01 00:00:00			
50%	4996.000000	2016-05-30 00:00:00			
75%	7494.000000	2017-04-09 00:00:00			
max	9994.000000	2017-12-31 00:00:00			
std	2884.906697	Nan			
	Ship Date	Postal Code	Sales Price	Quantity	\
count	9993	9993.000000	9993.000000	9975.000000	
mean	2016-04-22 21:14:17.039928064	55196.450515	229.897388	3.790376	
min	2014-01-07 00:00:00	1040.000000	0.444000	1.000000	
25%	2015-05-07 00:00:00	23223.000000	17.280000	2.000000	

```

50%           2016-06-12 00:00:00  56560.000000    54.528000   3.000000
75%           2017-04-29 00:00:00  90008.000000    209.970000   5.000000
max          2018-05-01 00:00:00  99301.000000   22638.480000  14.000000
std            NaN  32063.943256    623.269553   2.226107

      Discount     Profit Order_Year_From_ID Order_Year_From_Date
count  9993.000000  9993.000000        9993.000000        9993.000000
mean    0.156158    28.662956       2015.723006       2015.720004
min     0.000000  -6599.978000       2014.000000       1999.000000
25%    0.000000    1.731000       2015.000000       2015.000000
50%    0.200000    8.671000       2016.000000       2016.000000
75%    0.200000   29.364000       2017.000000       2017.000000
max    0.800000  8399.976000       2017.000000       2029.000000
std    0.206356   234.271064       1.124072        1.199837

```

[11]: df.isnull().sum()

```

[11]: Row ID          0
       Order ID        0
       Order Date       0
       Ship Date        0
       Ship Mode       98
       Customer ID      0
       Customer Name     0
       Segment          0
       Country          0
       City              0
       State             0
       Postal Code       0
       Region            0
       Product ID        0
       Category          0
       Sub-Category       0
       Product Name       0
       Sales Price        0
       Quantity          18
       Discount          0
       Profit             0
       Order_Year_From_ID 0
       Order_Year_From_Date 0
       dtype: int64

```

[12]: # # Imputation of Missing Values:
1. Impute missing values in the Ship Mode column using the calculated Days to
→Ship column.
2. Calculate Days to Ship as the difference between Ship Date and Order Date.
→If Days to Ship is 0,

```
# set Ship Mode to "Same Day"; if it is 7, set Ship Mode to "Standard Class"
# 3. Impute missing values in the Quantity column using a method of your choice.
    ↵ Print the rationale for selecting the method for imputation.
df['Days_to_Ship'] = (df['Ship Date'] - df['Order Date']).dt.days
```

[13]: # Fill missing Ship Mode values using Days_to_Ship logic

```
df.loc[(df['Ship Mode'].isna()) & (df['Days_to_Ship'] == 0), 'Ship Mode'] = ↵ 'Same Day'
df.loc[(df['Ship Mode'].isna()) & (df['Days_to_Ship'] == 7), 'Ship Mode'] = ↵ 'Standard Class'
df['Ship Mode'] = df['Ship Mode'].fillna(df['Ship Mode'].mode()[0])
```

[14]: # Impute missing Quantity using median (robust choice)

```
median_quantity = df['Quantity'].median()
df['Quantity'] = df['Quantity'].fillna(median_quantity)
```

[15]: df.isnull().sum()

```
Row ID          0
Order ID        0
Order Date      0
Ship Date        0
Ship Mode        0
Customer ID     0
Customer Name    0
Segment          0
Country          0
City              0
State             0
Postal Code       0
Region            0
Product ID        0
Category          0
Sub-Category      0
Product Name      0
Sales Price        0
Quantity          0
Discount          0
Profit            0
Order_Year_From_ID 0
Order_Year_From_Date 0
Days_to_Ship      0
dtype: int64
```

[16]: # 5. Data Masking and String Handling:

1. Drop the Customer Name column to protect Personal Identifiable Information ↵ (Pii).

```

# 2. Create a new column called Customer Name Masked, containing only the
# initials of the customer name.
df["Customer Name Masked"] = df["Customer Name"].apply(lambda x: ''.
    ↪join([name[0].upper() for name in x.split()]))
df.drop(columns=['Customer Name'], inplace=True)

# Convert the Postal Code column from numeric to text format, ensuring all
# codes are 5 characters long. Add a leading '0' where necessary.
df['Postal Code'] = df['Postal Code'].astype(str).str.zfill(5)

```

[17]: # 6. Data Type Conversion:

```

# 1. Convert the Quantity and Sales Price columns from strings to their
# appropriate numeric types (int and float, respectively).

df['Quantity'] = pd.to_numeric(df['Quantity'], errors='coerce').astype('Int64')
df['Sales Price'] = pd.to_numeric(df['Sales Price'], errors='coerce').
    ↪astype('float')

```

[18]: # 7. Handling Inconsistent Categorical Data:

```

# 1. Clean the State column by replacing abbreviations with full state names (e.
# ↪g., "CA" should be changed to "California").
# You may need to research state abbreviations online to ensure all entries are
# ↪corrected consistently.

```

```

print(df['State'].unique())
state_mapping = {
    'ca': 'california',
    'tx': 'texas',
    'ny': 'new york',
    'nj': 'new jersey',
    'wa': 'washington',
    'wa\\': 'washington'}

```

```

df["State"] = df["State"].replace(state_mapping)
df['State'] = df['State'].str.strip().str.title()

```

```

['kentucky' 'california' 'florida' 'north carolina' 'washington' 'texas'
 'wisconsin' 'utah' 'nebraska' 'pennsylvania' 'illinois' 'minnesota'
 'michigan' 'delaware' 'indiana' 'new york' 'arizona' 'virginia'
 'tennessee' 'tx' 'alabama' 'south carolina' 'oregon' 'colorado' 'iowa'
 'ohio' 'missouri' 'oklahoma' 'new mexico' 'louisiana' 'connecticut'
 'new jersey' 'massachusetts' 'georgia' 'nevada' 'rhode island'
 'mississippi' 'arkansas' 'montana' 'ca' 'new hampshire' 'maryland'
 'district of columbia' 'wa\\' 'nj' 'kansas' 'vermont' 'maine'
 'south dakota' 'idaho' 'north dakota' 'wyoming' 'west virginia' 'ny']

```

```
[19]: # 8. Feature Engineering:

# 1. Original Price: The price before any discount is applied.
df["Original Price"] = df["Sales Price"] / (1 - df["Discount"])

# 2. Total Sales: The total revenue generated by multiplying the Sales Price by
#   ↪Quantity.
df["Total Sales"] = df["Sales Price"] * df["Quantity"]

# 3. Total Profit: The total profit earned by multiplying the Profit by
#   ↪Quantity.
df["Total Profit"] = df["Profit"] * df["Quantity"]

# 4. Discount Price: The amount of discount applied, calculated based on the
#   ↪Original Price and Discount.
df["Discount Price"] = df["Original Price"] * df["Discount"]

# 5. Total Discount: The total discount value for the quantity sold.
df["Total Discount"] = df["Discount Price"] * df["Quantity"]
```

```
[20]: pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

```
[21]: df.head(5)
```

	Row ID	Order ID	Order Date	Ship Date	Ship Mode	Customer ID	\
0	1	ca-2016-152156	2016-08-11	2016-11-11	second class	cg-12520	
1	2	ca-2016-152156	2016-08-11	2016-11-11	second class	cg-12520	
2	3	ca-2016-138688	2016-12-06	2016-06-16	second class	dv-13045	
3	4	us-2015-108966	2015-11-10	2015-10-18	standard class	so-20335	
4	5	us-2015-108966	2015-11-10	2015-10-18	standard class	so-20335	

	Segment	Country	City	State	Postal Code	Region	\
0	consumer	united states	henderson	Kentucky	42420	south	
1	consumer	united states	henderson	Kentucky	42420	south	
2	corporate	united states	los angeles	California	90036	west	
3	consumer	united states	fort lauderdale	Florida	33311	south	
4	consumer	united states	fort lauderdale	Florida	33311	south	

	Product ID	Category	Sub-Category	\
0	fur-bo-10001798	furniture	bookcases	
1	fur-ch-10000454	furniture	chairs	
2	off-la-10000240	office supplies	labels	
3	fur-ta-10000577	furniture	tables	
4	off-st-10000760	office supplies	storage	

	Product Name	Sales Price	Quantity	\
--	--------------	-------------	----------	---

0	bush somerset collection bookcase	261.9600	2
1	hon deluxe fabric upholstered stacking chairs,...	731.9400	3
2	self-adhesive address labels for typewriters b...	14.6200	2
3	bretford cr4500 series slim rectangular table	957.5775	5
4	eldon fold 'n roll cart system	22.3680	2

	Discount	Profit	Order_Year_From_ID	Order_Year_From_Date	Days_to_Ship	\
0	0.00	41.9136		2016	2016	92
1	0.00	219.5820		2016	2016	92
2	0.00	6.8714		2016	2016	-173
3	0.45	-383.0310		2015	2015	-23
4	0.20	2.5164		2015	2015	-23

	Customer Name Masked	Original Price	Total Sales	Total Profit	\
0	CG	261.96	523.92	83.8272	
1	CG	731.94	2195.82	658.746	
2	DVH	14.62	29.24	13.7428	
3	SO	1741.05	4787.8875	-1915.155	
4	SO	27.96	44.736	5.0328	

	Discount	Price	Total Discount
0	0.0000	0.0	
1	0.0000	0.0	
2	0.0000	0.0	
3	783.4725	3917.3625	
4	5.5920	11.184	

[22]: # 2. Create a new column *Shipping Urgency* based on Days to Ship:

```
def categorize_urgency(days):
    if days == 0:
        return 'Immediate'
    elif 1 <= days <= 3:
        return 'Urgent'
    else:
        return 'Standard'

df['Shipping Urgency'] = df['Days_to_Ship'].apply(categorize_urgency)
```

[23]: #3. Create a column that calculates days since last order

```
df = df.sort_values(['Customer ID', 'Order Date'])
df['Days Since Last Order'] = df.groupby('Customer ID')['Order Date'].diff().dt.
    days
```

[24]: #4. Create a new dataset which stores the total sales, quantity and discount per customer and then merge these back to the original dataset

```
customer_summary = df.groupby('Customer ID').agg({
    'Total Sales': 'sum',
```

```

    'Quantity': 'sum',
    'Discount': 'mean'
}).reset_index()

customer_summary.rename(columns={
    'Total Sales': 'Customer Total Sales',
    'Quantity': 'Customer Total Quantity',
    'Discount': 'Customer Avg Discount'
}, inplace=True)

df = df.merge(customer_summary, on='Customer ID', how='left')

```

[25]: # Removing outlier

```

def remove_outliers(df, column):
    """
    Removes extreme outliers from a given column using the 3xIQR rule.
    Returns a cleaned DataFrame and prints the number of rows removed.
    """

    # Compute Q1 (25th percentile) and Q3 (75th percentile)
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1

    # Define bounds using 3*IQR
    lower_bound = Q1 - 3 * IQR
    upper_bound = Q3 + 3 * IQR

    # Track before removal
    initial_rows = df.shape[0]
    initial_order_ids = df['Order ID'].nunique()

    # Filter data within bounds
    df_clean = df[(df[column] >= lower_bound) & (df[column] <= upper_bound)]

    # Track after removal
    final_rows = df_clean.shape[0]
    final_order_ids = df_clean['Order ID'].nunique()

    print(f"Outlier removal in '{column}':")
    print(f"  Rows removed: {initial_rows - final_rows}")
    print(f"  Distinct Order IDs affected: {initial_order_ids - final_order_ids}")

    return df_clean

```

[26]: df = remove_outliers(df, 'Sales Price')

Outlier removal in 'Sales Price':
 Rows removed: 668
 Distinct Order IDs affected: 158

```
[27]: # 10. Customer Segmentation and Analysis:  

# 1. Calculate Customer Sales Quintile and Customer Profit Quintile based on  

# total sales and total profit per Customer ID.  

# 3. Create a cross-grid (cross-tabulation) based on these two quintiles to  

# analyze the relationship between customer sales and profitability.  

customer_summary = df.groupby('Customer ID').agg({'Total Sales': 'sum','Total  

#Profit': 'sum'}).reset_index()  

customer_summary['Customer Sales Quintile'] = pd.qcut(customer_summary['Total  

#Sales'],q=5,labels=['Q1', 'Q2', 'Q3', 'Q4', 'Q5'])  

customer_summary['Customer Profit Quintile'] = pd.qcut(customer_summary['Total  

#Profit'],q=5,labels=['Q1', 'Q2', 'Q3', 'Q4', 'Q5'])  

cross_tab = pd.crosstab(customer_summary['Customer Sales  

#Quintile'],customer_summary['Customer Profit Quintile'],normalize='index')  

print(cross_tab)
```

Customer Profit Quintile	Q1	Q2	Q3	Q4	Q5
Customer Sales Quintile					
Q1	0.245283	0.547170	0.201258	0.006289	0.000000
Q2	0.196203	0.215190	0.398734	0.189873	0.000000
Q3	0.145570	0.120253	0.202532	0.373418	0.158228
Q4	0.215190	0.063291	0.120253	0.272152	0.329114
Q5	0.201258	0.050314	0.075472	0.157233	0.515723

insights:- Strong Positive Relationship between Sales and Profit: The Q5–Q5 shows that around 51.6% of the highest sales customers are also in the highest profit group. that means that our top spenders are also your most profitable customers.

Low Sales → Low Profit: Most Q1 (Low Sales) customers fall in Q1–Q2 (Low Profit) ranges — over 79% (0.245 + 0.547). This shows that customers who buy less also contribute little to profit

Middle Segments Show Mix: Quintiles Q2–Q3 are spread across profit quintiles — they include both low- and mid-profit customers. This indicates a mixed performance group

High Sales but Medium Profit (Q5–Q3 or Q5–Q4): A smaller portion around 23% of high sales customers fall in Q3–Q4 profit quintiles. These could be discount-driven buyers — high revenue but lower margins.

Business Insight: Focus retention efforts on Q5–Q5 customers — they drive both revenue and profit. Reassess pricing or discount strategies for high-sales but medium-profit customers. Identify ways to move Q3 and Q4 sales customers up into higher profit brackets by offering premium products or optimizing shipping/discount policies.

[]:

Final Dashboard:-


```
[28]: # 1. Sales and Profit Analysis:

# 1. Top 10 Most Profitable Products: Use a bar chart to display the products with the highest total profit
# 2. Top 10 Most Loss Making Products: Use a bar chart to display the products with the highest total losses (negative profit).
# 3. Sales vs. Profit Correlation: Use a scatter plot to visualize the correlation between Total Sales and Total Profit.
# Add a regression line to show the trend.
```

```
[29]: import seaborn as sns
import matplotlib.pyplot as plt

# Top 10 most profitable products
top_profitable = (df.groupby('Product Name')['Total Profit'].sum().
                   sort_values(ascending=False).head(10))


plt.figure(figsize=(10, 5))
sns.barplot(x=top_profitable.values, y=top_profitable.index,hue = top_profitable.index, palette='Greens_r', legend = False)
plt.title("Top 10 Most Profitable Products")
plt.xlabel("Total Profit")
plt.ylabel("Product Name")
plt.show()
```



These products contribute significantly to overall profit. Most of them belong to office supplies and stationery categories such as labels, binders, and papers. They have high sales volume and good profit margins, indicating steady customer demand and efficient pricing.


Actions:- They should be prioritized for promotion, inventory stocking, and marketing focus

```
[30]: top_loss = (df.groupby('Product Name')['Total Profit'].sum() .
    ↪sort_values(ascending=True).head(10))
plt.figure(figsize=(10, 5))
sns.barplot(x=top_loss.values, y=top_loss.index,hue = top_loss.index,
    ↪palette='Reds_r', legend = False)
plt.title("Top 10 Most Loss-Making Products")
plt.xlabel("Total Loss")
plt.ylabel("Product Name")
plt.show()
```



These products generate negative profits every sale results in a loss. Many are office equipment and technology-related items like printers, binding machines, and box systems, which usually have higher costs or are sold under heavy discounts. Heavy discounts or promotional sales. High shipping or handling costs (especially for large/heavy products). Returns or warranty replacements increasing cost. Low demand, leading to inventory holding cost or clearance sales

```
[31]: plt.figure(figsize=(8,6))
sns.scatterplot(data=df, x="Total Sales", y="Total Profit", alpha=0.6)
sns.regplot(data=df, x="Total Sales", y="Total Profit", scatter=False,
    ↪color="red")
plt.title("Sales vs Profit Correlation")
plt.xlabel("Sales Price")
plt.ylabel("Total Profit")
plt.tight_layout()
plt.show()
```



```
[32]: correlation = df[\"Total Sales\"].corr(df[\"Total Profit\"])\nprint(f\"Correlation between Sales and Profit: {correlation:.2f}\")
```

Correlation between Sales and Profit: 0.23

The scatter plot shows a slightly positive correlation between Sales and Profit — meaning that generally, higher sales tend to bring higher profit. However, the relationship is not very strong (the points are widely spread). You can notice that some orders with very high sales still have low or negative profit. This usually happens when high discounts or shipping costs reduce profit margins. The red regression line slopes upward but slightly — indicating a weak positive trend between total sales and total profit. There are some points far below the line (negative profit despite good sales).

These represent loss-making transactions, likely due to: Heavy discounts High-cost shipping modes Returns or refunds

```
[34]: # 3. Shipping and Delivery Analysis:
```


- # 1. Distribution of Shipping Urgency: Visualize the distribution of orders by ↵ Shipping Urgency using a pie chart or bar chart
- # 2. Days to Ship vs. Profit: Use a box plot to explore the distribution of ↵ Profit across different Days to Ship categories.

```

# This will help analyze whether faster shipping correlates with higher or lower profitability.
# 3. Shipping Mode and Profitability: Create a grouped bar chart to compare the profitability of different shipping modes
# (eg.. Standard Class,First Class).
# 4. Using pivot table, determine which shipping modes are most preferred across different regions and analyze the impact on total sales and profit.
# Create a pivot table that shows the count of Order IDs, total Sales, and total Profit for each Region and Ship Mode. Identify and print your insights

```

```
[60]: df['Shipping Urgency'].value_counts().plot(kind='pie', autopct='%.1f%%', startangle=90, colors=['#66c2a5', '#fc8d62', '#8da0cb'])
plt.ylabel('')
plt.title('Shipping Urgency Distribution')
plt.show()
```



The data shows that Standard Shipping accounts for ~80% of all orders, while Urgent orders make up 15.1%, and Immediate (Same Day) orders only 4.9%. This indicates that most customers prefer cost-effective standard delivery over faster options. The low percentage of Immediate orders suggests that customers are not highly time-sensitive and are likely price-conscious. From a business perspective, this pattern highlights an opportunity to optimize logistics around standard shipping, as it handles the majority of orders. The company can also limit high-cost same-day options to

select products or regions to reduce operational expenses.


```
[61]: plt.figure(figsize=(8,5))
sns.boxplot(data=df, x='Shipping Urgency', y='Total Profit', color='skyblue')
plt.title('Distribution of Profit Across Shipping Urgency')
plt.xlabel('Shipping Urgency')
plt.ylabel('Total Profit')
plt.show()
```



The primary insight is the striking similarity across all three shipping urgency levels.: Since neither Urgent nor Immediate shipping yields a higher profit than Standard shipping, the company should focus on minimizing the use of these faster, more expensive methods to reduce operational costs without sacrificing transactional profit.

```
[65]: shipping_profit = df.groupby('Ship Mode')['Profit'].sum().
    sort_values(ascending=False).reset_index()

plt.figure(figsize=(8,5))
sns.barplot(data=shipping_profit, x='Ship Mode', y='Profit')
plt.title('Total Profit by Shipping Mode')
plt.xlabel('Ship Mode')
plt.ylabel('Total Profit')
plt.show()
```



Key Insights

- Standard Class Dominance:** The ‘Standard Class’ ship mode is the largest contributor to total profit, generating well over \$70,000 in aggregate profit.
- Second Class Contribution:** ‘Second Class’ is the second-highest contributor, generating around \$28,000 in total profit.
- First Class and Same Day** significantly less.

Focus on Standard: The company’s overall profitability is heavily reliant on orders shipped via ‘Standard Class.’ This shipping mode is the backbone of the company’s total earnings. the company is likely better off by limiting high-cost/fast shipping options

```
[66]: pivot = pd.pivot_table(df,index="Region",columns="Ship Mode",values=["Order ID", "Total Sales", "Total Profit"],
                            aggfunc={"Order ID": "nunique", "Total Sales": "sum", "Total Profit": "sum"},
                            fill_value=0,
                            margins=True,
                            margins_name="Total")
pivot
```

```
[66]:          Order ID
           Ship Mode first class same day second class standard class   Total \
Region
central             168        60       220        692    1140.0
east               223        69       254        808    1354.0
```

south	122	37	155	482	796.0
west	247	86	304	924	1561.0
Total	760.0	252.0	933.0	2906.0	4851.0

Total Profit		\			
Ship Mode	first class	same day	second class	standard class	Total
Region					
central	379.5891	3234.023	8464.0572	-2857.5172	9220.1521
east	21950.2402	7019.1758	25494.5343	104570.3065	159034.2568
south	16070.5882	7933.404	34050.1345	78352.9964	136407.1231
west	23741.0491	11673.3681	60498.3407	134145.9009	230058.6588
Total	62141.4666	29859.9709	128507.0667	314211.6866	534720.1908

Total Sales		\			
Ship Mode	first class	same day	second class	standard class	Total
Region					
central	131216.004	49495.856	233994.8	658172.199	1072878.859
east	209504.811	67024.471	252083.511	758728.405	1287341.198
south	103478.698	38692.503	164570.0975	462073.992	768815.2905
west	247706.9135	79858.384	323473.3055	945701.568	1596740.171
Total	691906.4265	235071.214	974121.714	2824676.164	4725775.5185

The West region generates the highest total profit (\$230,058.66), followed by the East region (\$159,034.26). Central and South Profit:Central (\$92,200.15) and South (\$78,352.99) lag significantly behind the West and East regions in terms of aggregate profit.

Reinforce the West and East: These two regions are your profit powerhouses (especially the West). Ensure operational efficiency is optimized here to maintain high sales and profit. Focus Improvement on Central and South:

[42]: # 4. Regional Sales and Profitability:

```
# 1. Sales and Profit by Region: Use a map or bar chart to visualize total sales and profit by region or state.
# This will highlight which regions are the most profitable
# 2. State-wise Profitability: Create a pivot table to summarize the profitability of each state. Highlight the top and bottom states based on profitability.
```

[43]: region_summary = df.groupby('Region')[['Total Sales', 'Total Profit']].sum().reset_index()

```
# Plot
plt.figure(figsize=(8,5))
sns.barplot(data=region_summary, x='Region', y='Total Sales', color='skyblue', label='Sales')
sns.barplot(data=region_summary, x='Region', y='Total Profit', color='orange', label='Total Profit')
```

```

plt.title('Total Sales and Profit by Region')
plt.ylabel('Value')
plt.legend()
plt.show()

```



Sales:- The West and East regions are the clear leaders in terms of Total Sales, with both regions generating sales exceeding \$1.2 million. Central is the next largest contributor, followed by South, which has the lowest sales volume (below \$0.8 million).

Profit:- The West region generates the highest Total Profit. The East region is the second highest profit contributor. Central and South regions contribute significantly less to the total profit.

Sales/profit:- The West region shows a healthy profit margin. Although its sales bar is the highest, the orange profit slice is proportionally very large relative to the blue sales slice. The East region, while having high sales, appears to have a slightly lower profit margin compared to the West. The South and East region's profit appears to be the thinnest relative to its sales bar, suggesting it has the lowest profit margin.

```

[44]: # Pivot table for State-wise total profit
state_profit = df.groupby('State')[['Profit']].sum().reset_index() .
    ↪sort_values(by='Profit', ascending=False)

print("Top 10 Most Profitable States:")

```

```

print(state_profit.head(10))

print("\nBottom 10 Least Profitable States:")
print(state_profit.tail(10))

```

Top 10 Most Profitable States:

	State	Profit
3	California	44672.2302
30	New York	28135.2368
45	Washington	10376.6167
20	Michigan	9824.2301
44	Virginia	7279.8420
9	Georgia	6946.9138
15	Kentucky	5335.4919
12	Indiana	5002.5213
28	New Jersey	4495.7085
19	Massachusetts	4470.3288

Bottom 10 Least Profitable States:

	State	Profit
7	District Of Columbia	94.1444
40	Tennessee	-101.0933
8	Florida	-696.0981
35	Oregon	-1414.3545
1	Arizona	-2135.8807
4	Colorado	-3438.5615
33	Ohio	-5339.8993
36	Pennsylvania	-9324.1787
11	Illinois	-11818.3018
41	Texas	-13031.3552

[45]: # 5. Discount and Pricing Analysis:

```

# 1. Impact of Discounts on Profitability: Use a scatter plot with a trend line
# to analyze how different levels of discount affect profitability.
# 2. Original Price vs. Discounted Price: Create chart to compare the original
# price and the discounted price across various product
# categories or sub-categories

```

[46]:

```

import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(8,6))
sns.scatterplot(data=df, x="Discount", y="Profit", alpha=0.6)
sns.regplot(data=df, x="Discount", y="Profit", scatter=False, color="red")
plt.title("Impact of Discount on Profitability")
plt.xlabel("Discount (%)")

```

```

plt.ylabel("Profit")
plt.tight_layout()
plt.show()

```



There is a strong negative correlation between the discount percentage and Profit. As the discount increases, the profit tends to decrease. At 0% discount, the profit is clustered tightly around zero, with some instances of high profit 400 and some moderate losses. As the discount approaches 30% to 40% the bulk of transactions start to fall below the zero profit line, meaning the orders are generally resulting in a loss. High discounts, particularly those around 60% to 80%, are where the most severe losses occur. The largest losses, dipping to -1,000 to -1,200, are concentrated at the highest discount levels 70% to 80%. This directly addresses the “negative outliers” problem seen in previous analyses—high discounts are a major driver of large losses.

Discount Strategy Review: The company must urgently review its discounting policy. Discounts, particularly those exceeding 40%, are highly likely to result in losses and are the main cause of the extreme financial outliers. Set a Discount Cap: A policy should be implemented to strictly control or eliminate discounts above a certain threshold unless the underlying product has an exceptionally high margin to absorb the markdown.

```
[35]: import numpy as np
import matplotlib.pyplot as plt
```

```


# Prepare data
price_comparison = df.groupby('Sub-Category')[['Original Price', 'Discount Price']].mean().reset_index().sort_values(by='Original Price')

# Set positions and width
x = np.arange(len(price_comparison['Sub-Category'])) # the label locations
width = 0.4 # the width of the bars

# Plot
plt.figure(figsize=(12,6))
plt.bar(x - width/2, price_comparison['Original Price'], width, label='Original Price', color='skyblue')
plt.bar(x + width/2, price_comparison['Discount Price'], width, label='Discount Price', color='orange')

# Labels and formatting
plt.xticks(x, price_comparison['Sub-Category'], rotation=45)
plt.ylabel('Average Price')
plt.title('Original vs Discounted Price by Sub-Category')
plt.legend()
plt.tight_layout()
plt.show()

```



Technology and Furniture Drive Dollar Losses: The two sub-categories with the largest average dollar discounts are Machines and Copiers, both from the Technology category. Furniture's Triple Threat (Tables, Bookcases, Chairs): The next three largest dollar discounts come from the Furniture category: Furniture has the lowest profit margin (2.49%)


Profit Killer (Furniture): The high dollar discounts on Tables, Bookcases, and Chairs part of the Furniture category, which has the lowest overall margin must be strictly controlled to prevent large individual transaction losses.

Revenue Killer (Technology): While Machines and Copiers can absorb large dollar discounts due to their high price, their profitability should be closely monitored to ensure the huge revenue they generate is not being eroded by discounts.

```
[47]: # Group and prepare the data
pf = df.groupby('Category')[['Original Price', 'Discount Price']].mean()
    ↪reset_index()

# Melt the DataFrame to long format for grouped bar chart
pf_melted = pf.melt(id_vars='Category', value_vars=['Original Price', 'Discount Price'],
    ↪var_name='Price Type', value_name='Average Price')

# Plotting
plt.figure(figsize=(10,6))
sns.barplot(data=pf_melted, x='Category', y='Average Price', hue='Price Type')
plt.title('Original vs Discounted Price by Category')
plt.ylabel('Average Price')
plt.xticks(rotation=45)
plt.legend(title='Price Type')
plt.tight_layout()
plt.show()
```



- Technology and Furniture categories have the highest original prices (200–250 range) and also show noticeable discounts, suggesting retailers offer more discounts on expensive items to boost purchase intent.
- Office Supplies have the lowest prices (< 100) and minimal discounts, meaning these are low-cost essentials where discounts aren't needed to drive sales.
- The gap between original and discounted prices is largest for Technology, indicating a high-discount strategy — likely to stay competitive or clear inventory of fast-evolving tech products.
- Furniture discounts are moderate — possibly due to higher shipping or storage costs that limit heavy markdowns.

```
[45]: df.groupby(df["Category"])["Total Profit"].sum().reset_index()
```

```
[45]:      Category  Total Profit
0    furniture     -2408.0523
1  office supplies    346667.4241
2    technology     190460.819
```

```
[49]: # 6. Temporal Analysis:
```

```
# 1. Sales and Profit Trends Over Time: Use a time series plot to analyze how sales and profit have trended over the years or months.
# This will help in identifying any seasonal patterns.
# 2. Order Frequency by Month: Use a bar chart or line plot to show the number of orders placed each month. Highlight any months with unusually high or low order frequencies.
# 3. Yearly Growth in Sales and Profit: Use a year-over-year growth chart to compare the sales and profit growth over different years.
```

```
[52]: monthly_trend
```

```
[52]:   YearMonth  Sales  Price  Profit
0    2014-01  15492.5080  1800.6074
1    2014-02   7376.9780  1161.5573
2    2014-03  15570.3810  1152.7621
3    2014-04  14806.4460  1810.4904
4    2014-05  18301.6400  2146.7556
5    2014-06  15775.6216  2058.4346
6    2014-07  18195.1120   669.9275
7    2014-08  16371.6835  1728.5532
8    2014-09  21843.8998  2774.2497
9    2014-10  19558.1110  2152.2706
10   2014-11  35540.6235  4551.1441
11   2014-12  32474.5407  2730.3293
12   2015-01   8050.2060  1046.9356
13   2015-02  13338.2100  1245.4814
14   2015-03  15166.7616  2493.5004
15   2015-04  18275.9125  2321.2967
16   2015-05  18725.1845   874.1701
```

17	2015-06	15591.4352	1858.9962
18	2015-07	18168.1210	1188.8596
19	2015-08	22443.2760	2454.1427
20	2015-09	30337.6750	5247.9759
21	2015-10	18714.7930	2770.0687
22	2015-11	37733.8835	4434.2015
23	2015-12	29678.7872	3010.8394
24	2016-01	13827.2200	1552.6127
25	2016-02	14987.4820	1720.7262
26	2016-03	25617.5730	381.1710
27	2016-04	17218.7650	1584.8303
28	2016-05	28830.5100	5012.9304
29	2016-06	23854.6890	2277.9347
30	2016-07	23372.9990	3646.4781
31	2016-08	25375.6508	1977.3070
32	2016-09	28882.0535	3195.7547
33	2016-10	26419.5520	3055.1815
34	2016-11	30697.7280	2697.2984
35	2016-12	28722.0905	2365.4004
36	2017-01	23669.9140	3362.4312
37	2017-02	23562.3874	260.5841
38	2017-03	35249.1218	4818.3228
39	2017-04	19834.6446	917.1285
40	2017-05	28075.5084	4310.3243
41	2017-06	29623.0120	3605.9494
42	2017-07	31376.2125	3799.5514
43	2017-08	27623.9290	4578.4098
44	2017-09	43481.8730	5646.2500
45	2017-10	27925.4382	146.0604
46	2017-11	39370.8280	3289.8976
47	2017-12	35134.0738	3553.5435

```
[51]: df[\"YearMonth\"] = df[\"Order Date\"].dt.to_period(\"M\").astype(str)
monthly_trend = df.groupby(\"YearMonth\")[[\"Sales Price\", \"Profit\"]].sum().
    reset_index()

import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(14,6))
sns.lineplot(data=monthly_trend, x=\"YearMonth\", y=\"Sales Price\", label=\"Sales\", marker=\"o\")
sns.lineplot(data=monthly_trend, x=\"YearMonth\", y=\"Profit\", label=\"Profit\", marker=\"o\")

plt.title(\"Monthly Sales and Profit Trends Over Time\")
plt.xticks(rotation=45)
```

```

plt.ylabel("Amount")
plt.xlabel("Year-Month")
plt.legend()
plt.tight_layout()
plt.show()

```



Sales show an upward trend over time, though with noticeable seasonal fluctuations. Profit follows a similar pattern but stays much lower than sales, indicating thin profit margins or high costs/discounts in some months.

There are repeated sales spikes toward year-end (Nov–Dec) — possibly due to holiday or festival seasons, when promotions and discounts increase demand. Profit does not always peak with sales, suggesting discount-heavy sales periods reduce profitability November and September consistently show high sales and profit, suggesting strong seasonal demand. February and July often have lower performance, possibly due to off-season effects or market slowdowns.

Some months show high sales but low profits or even dips, meaning operational inefficiencies or over-discounting could be affecting profitability March 2016: High sales (25,617.57) but very low profit (381.17), indicating potential discounting, high costs, or operational inefficiencies. October 2017: Decent sales (27,925.44) but lowest profit (146.06), possibly due to unexpected expenses or pricing issues.

The business is growing in sales, but profit optimization is needed. Focus should be on reducing costs or optimizing discount strategies during high-sales months to improve profit margins.


```

[59]: plt.figure(figsize=(12,6))
sns.barplot(data=monthly_orders,x="Month",y="Order_U
↪Count",hue="Month",palette="coolwarm",legend=False)

plt.title("Order Frequency by Month")
plt.ylabel("Number of Orders")
plt.xticks(rotation=45)

```

```
plt.tight_layout()
plt.show()
```



Peak Ordering Months: The highest number of orders occur in the last quarter (October, November, December). This confirms a strong year-end sales surge, likely due to festive seasons, holiday shopping, and promotional offers.

Moderate Activity: Mid-year months (April–August) show average order volume — business remains steady but not at its highest. This may reflect regular business demand without seasonal effects.

Low Order Period January–February record the lowest order counts, suggesting a post-holiday slowdown where customers and businesses reduce spending after major festivals or holidays.

The data shows a clear upward trend from mid-year to year-end, showing seasonality in customer purchase behavior. The company experiences strong Q4 performance, so marketing and inventory should focus on this period. Plan stock replenishment, staffing, and marketing campaigns ahead of Q4 to meet high demand. Consider boosting sales in slow months (like Jan–Feb) with discounts, loyalty programs, or new product launches.

```
[54]: df["Year"] = df["Order Date"].dt.year
yearly_trend = df.groupby("Year")[["Sales Price", "Profit"]].sum().reset_index()
yearly_trend["Sales Growth (%)"] = yearly_trend["Sales Price"].pct_change() * 100
yearly_trend["Profit Growth (%)"] = yearly_trend["Profit"].pct_change() * 100

import seaborn as sns
import matplotlib.pyplot as plt


plt.figure(figsize=(12,6))
```

```

sns.lineplot(data=yearly_trend, x="Year", y="Sales Growth (%)", marker="o", u
↳label="Sales Growth")
sns.lineplot(data=yearly_trend, x="Year", y="Profit Growth (%)", marker="o", u
↳label="Profit Growth")

plt.title("Year-over-Year Growth in Sales and Profit")
plt.ylabel("Growth (%)")
plt.xticks(yearly_trend["Year"])
plt.axhline(0, color='gray', linestyle='--')
plt.legend()
plt.tight_layout()
plt.show()

```



2014–2015: Stable & Profitable Growth Both sales and profit increased, but profit grew faster (17%) than sales (6%). This means better efficiency or higher margins — the business was scaling profitably.

2015–2016: Warning Signal Sales rose by 16.9%, but profit almost stagnated (only 1.8%). This shows that despite higher revenue, costs increased sharply or low-margin products dominated sales.

2016–2017: Recovery & Strong Growth Both sales (+26.8%) and profit (+29.9%) grew strongly. This indicates strategic improvements likely better cost control, reduced discounts, or focus on high-profit items. Business became more efficient and sustainable again.

[]: